



COURSE SLIDES

Agentic AI for LinkedIn Marketing (C386)

Course Code: C386

Conducted by Tertiary Infotech Academy Pte Ltd · UEN 201200696W

Version v1.0 · 3 August 2026

© 2026 Tertiary Infotech Academy Pte Ltd. All rights reserved. · www.tertiarycourses.com.sg



COURSE ADMINISTRATION

Welcome & Housekeeping

About the Trainer



Your Trainer

General Trainer template —
to be completed by the trainer

NAME

TITLE / DESIGNATION

QUALIFICATIONS

AREAS OF EXPERTISE

TRAINING & INDUSTRY EXPERIENCE

CONTACT

About the Trainer



Tertiary Infotech Academy

Principal Trainer
Tertiary Infotech Academy Pte. Ltd.

ROLE

Principal Trainer, Tertiary Infotech Academy Pte. Ltd.

CERTIFICATION

Industry certified in the subject area of this course.

DELIVERS

Professional short courses for Tertiary Infotech Academy.

FOUNDER

Founder and lead instructor at Tertiary Infotech / Tertiary Courses.

Let's Know Each Other

1 Your name and organisation / role.

2 Your experience with this subject (if any).

3 What you want to be able to do after this course.

Ground Rules

1

Set your mobile phone to silent mode.

2

Participate actively — no question is too small.

3

Mutual respect: agree to disagree.

4

One conversation at a time.

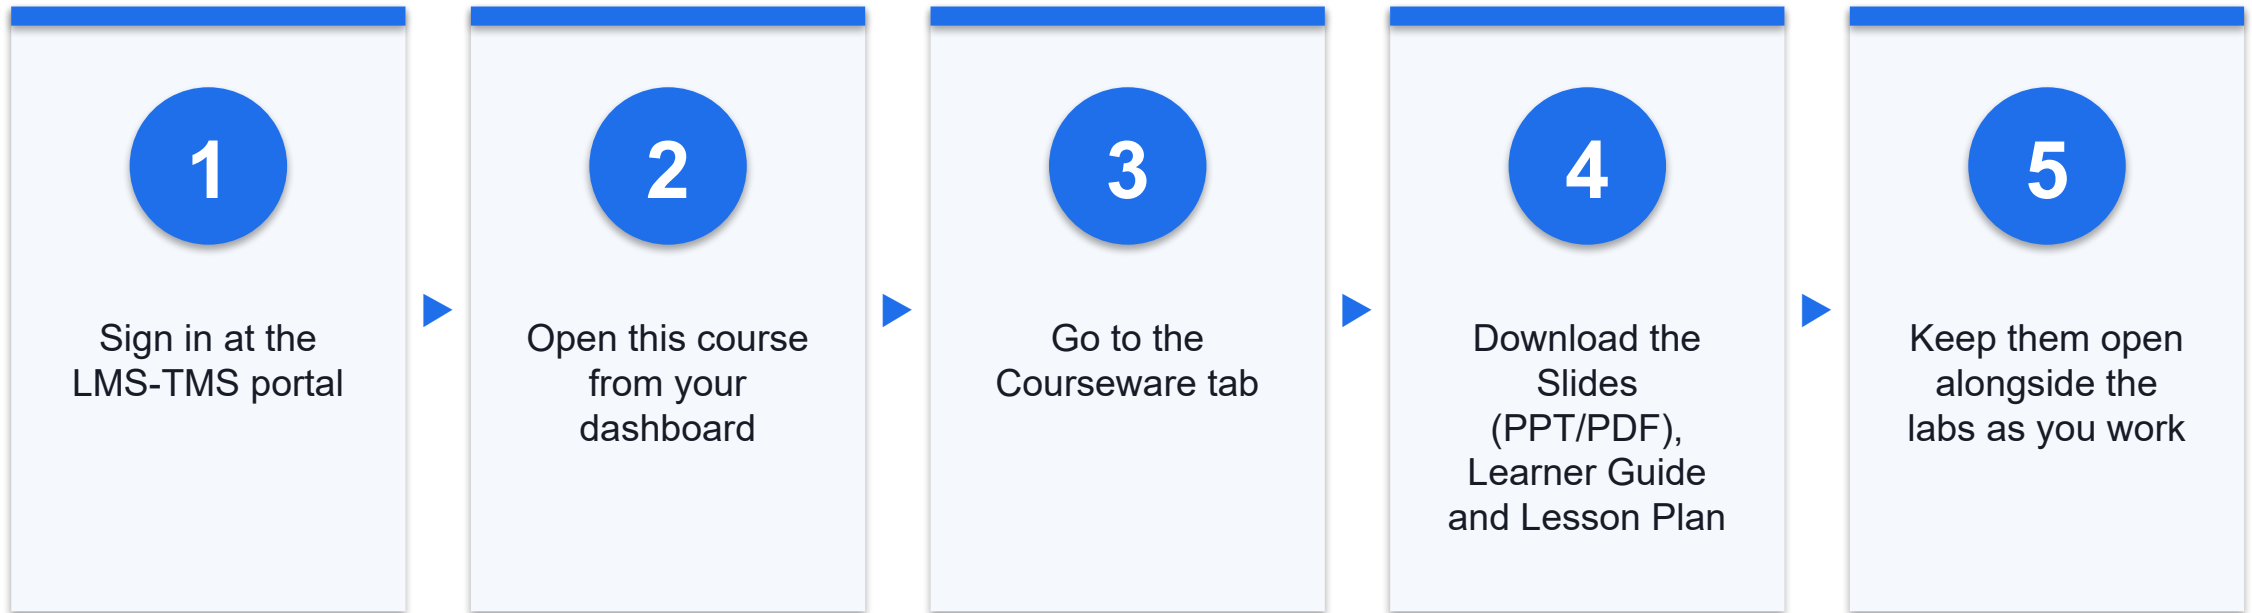
5

Be punctual; return from breaks on time.

6

Learn by doing — try every hands-on lab.

Download Course Material



Portal: <https://lms-tms.tertiaryinfotech.com>

Lesson Plan — 2 Days, 8 hours/day

Day 1

- Day 1 — Agent foundations, profile, prompting, and content systems
 - Topic 1: Getting Started with Agentic AI for LinkedIn Marketing (Labs 1–2)
 - Topic 2: Content Creation and Personal Branding with AI (Labs 3–4)

Day 2

- Day 2 — Prospecting, outreach, governed workflows, and improvement
 - Topic 3: Lead Generation and Outreach with AI Agents (Labs 5–6)
 - Topic 4: Automating and Scaling LinkedIn with AI Agents (Labs 7–8)
- Daily timing
 - 9:00 am - 6:00 pm (1-hour lunch; two 15-minute tea breaks)

Learning Outcomes

1

Agent Foundations

Explain how goal-driven AI agents support LinkedIn marketing and define safe tool, evidence, and human-approval boundaries.

2

Profile & Prompts

Optimise a LinkedIn profile and write reusable prompts that preserve an authentic professional voice.

3

Content System

Create evidence-grounded LinkedIn posts, articles, visual briefs, carousels, and a coherent content calendar.

4

Prospecting

Research and segment prospects from authorised information, then draft relevant connection and nurture messages for human review.

5

Governed Workflow


Design a draft-first content, engagement, and follow-up workflow that respects LinkedIn rules and professional etiquette.

6

Measure & Improve

Define LinkedIn performance metrics, diagnose synthetic results, and recommend one bounded improvement with a traceable decision record.

How You'll Learn



CORE CONCEPTS

Agentic LinkedIn Marketing Fundamentals

Four Ideas to Keep in View

1

Strategy

Tie profile, content, outreach, and measurement to one decision chain.

2

Evidence

Keep approved facts, labelled hypotheses, and unknowns distinct.

3

Human authority

External actions stay behind named approval gates.

4

Learning loop

Validate, diagnose, propose one change, observe, and record.

The G-C-A-T-E Agent Contract

1

Goal

Define the decision and useful outcome.

2

Context & Audience

Supply approved sources and reader needs.

3

Task & Evidence

Specify the output and traceability rules.

4

Evaluation

Set checks, permissions, stop conditions, and handoff.

OPERATING PRINCIPLE

Automate preparation; preserve professional judgement.

The agent researches approved inputs, drafts, checks, and recommends. A person owns identity, relationships, publication, messaging, and correction.

What You Will Build

Brand & Content

- Profile value stack
- Voice fingerprint
- Content kit and calendar

Prospect & Outreach

- Evidence-based segments
- Personalised draft sequence
- Etiquette and escalation

Operate & Improve

- Draft-first workflow
- Approval packet
- Metric-led improvement loop

The Connected Learning Arc

1

Frame the goal, profile, evidence, tools, and authority boundary.

2

Turn those decisions into a tested reusable prompt contract.

3

Create content and a calendar using the same voice and evidence.

4

Segment synthetic prospects and draft respectful outreach.

5

Design a governed workflow, measure results, and recommend one bounded improvement.

TOPIC 01

Getting Started with Agentic AI for LinkedIn Marketing

LinkedIn strategy | AI agents and tools | profile and personal brand | effective prompting

KEY CONCEPTS

Key Concepts — Getting Started with Agentic AI for LinkedIn Marketing

1

Goal before tool

Begin with the business result, audience need, desired professional action, and evidence of success.

2

Agent loop

A bounded agent observes approved context, reasons against instructions, drafts an output, checks it, and hands off.

3

Model, tools, instructions

The model interprets; tools retrieve or act; instructions define scope, quality, and stop conditions.

4

Human authority

Profile edits, public posts, invitations, messages, and account changes remain explicit human decisions.

5

Professional identity

A useful profile connects credible experience, audience value, proof, and a clear next action without keyword stuffing.

6

Evidence boundary

Approved facts, labelled hypotheses, and unknowns are kept distinct so fluent output is not mistaken for truth.

Key Concepts — Getting Started with Agentic AI for LinkedIn Marketing (cont.)

7

Prompt contract

Goal, Context, Audience, Task, Evidence, and Evaluation turn a vague request into a reusable operating instruction.

8

Responsible use

No scraping, unauthorised bots, mass messaging, credential sharing, deceptive claims, or inauthentic engagement.

LinkedIn Marketing as a Decision System


1

LinkedIn activity is useful only when it supports a defined professional or business result. A decision chain connects the result to an audience need, a desired action, a channel job, and a metric. This prevents an agent from optimising for visible activity while ignoring whether the right people are learning, trusting, or taking the intended next step.

2

For example, a consultancy seeking qualified discovery calls may use profile visits and relevant replies as leading indicators, but it should not call an impression a lead. The profile, content, and outreach system must all point to the same value proposition and next action.

LinkedIn Marketing as a Decision System — Concept Map



What Makes a Marketing Workflow Agentic

1

A single AI-generated post is assistance, not an agent. An agentic workflow manages a multi-step goal: retrieve authorised context, choose a routine, produce a draft, evaluate it against rules, request missing information, and stop or hand control back when risk is high.

2

The safest beginner pattern is recommendation-only. The agent may read approved synthetic files, draft artefacts, and recommend actions. It cannot publish, invite, message, collect personal data, or change account settings. This design is useful because it improves repeatability without giving away human accountability.

What Makes a Marketing Workflow Agentic — Concept Map

1

Observe

Read approved brief, profile, signals, or results.

2

Reason

Apply goal, evidence rules, and decision criteria.

3

Draft

Create a structured recommendation or content draft.

4

Check & hand off

Validate, log, and pause for the authorised person.

Profile and Personal-Brand Value Stack


1

The introduction section is the first profile area many visitors see. A strong headline states credible role or capability, the audience or problem served, and a differentiating proof point. The About section expands this into a concise story: current focus, audience value, evidence, point of view, and invitation to continue the conversation.

2

Optimisation is not keyword repetition. LinkedIn advises members to keep profiles rich, accurate, and complete and to avoid keyword stuffing, inflated experience, and using a personal profile as an advertisement. Every AI edit therefore needs a source note and an accuracy check by the profile owner.

Profile and Personal-Brand Value Stack — Concept Map



The G-C-A-T-E Prompt Contract

1

A prompt contract is a reusable instruction set, not a one-off clever phrase. G-C-A-T-E names the Goal, Context, Audience, Task, Evidence, and Evaluation requirements. It also defines allowed tools, output shape, human gates, and failure behaviour.

2

A good contract makes missing information visible. If a proof point is absent, the agent writes UNKNOWN or asks for it; it does not invent a client result. Evaluation criteria should include factual support, professional relevance, voice match, platform fit, and a clear next action.

The G-C-A-T-E Prompt Contract — Concept Map

1

G - Goal

Name the decision or useful outcome.

2

C/A - Context & Audience

Supply approved facts and reader needs.

3

T/E - Task & Evidence

Specify output and traceable inputs.

4

E - Evaluation

Define checks, limits, and stop conditions.

Risk Tiers and Approval Gates


1

Risk increases when a workflow moves from synthetic analysis to external action. Reading a supplied brand brief is low risk; drafting a profile or message is medium risk; publishing, inviting, messaging, processing personal data, or changing account settings is high risk.

2

An approval gate states the action, owner, evidence packet, decision options, and rollback. A useful status vocabulary is DRAFT, READY FOR HUMAN REVIEW, APPROVED BY ROLE, and STOP - REASON. The agent must never treat silence as approval.

Risk Tiers and Approval Gates — Concept Map



Hands-On Labs — Getting Started with Agentic AI for LinkedIn Marketing

Lab 1

- Build the LinkedIn Strategy, Profile, and Agent Boundary Map

Lab 2

- Write and Boundary-Test the LinkedIn G-C-A-T-E Prompt Contract

Build the LinkedIn Strategy, Profile, and Agent Boundary Map

LAB 1

You start the connected Northstar Advisory scenario by turning its approved synthetic brand brief, audience signals, and current-profile draft into one coherent LinkedIn foundation. You identify what an AI agent may read, draft, recommend, or never do, then map the profile value stack and the decisions a human must retain.

You'll build

C386-linkedln-pack/01-foundation-profile-and-boundaries.md containing the decision chain, audience needs, profile audit, value stack, tool map, risk tiers, and approval gates.

Tools: Approved AI assistant, text editor, supplied Markdown and CSV resources

Build the LinkedIn Strategy, Profile, and Agent Boundary Map

STEP 1 OF 2



Frame the decision chain and profile value stack from the approved brand brief, audience signals, and current-profile draft.

\$ FACT / HYPOTHESIS / UNKNOWN

Build the LinkedIn Strategy, Profile, and Agent Boundary Map

STEP 2 OF 2

2

Map tools, risk tiers, and explicit human gates before any identity, publication, invitation, message, data, or account action.

\$ DRAFT -> HUMAN REVIEW -> APPROVE / REVISE / STOP

Build the LinkedIn Strategy, Profile, and Agent Boundary Map

Test it

Open 01-foundation-profile-and-boundaries.md. It must contain five audience-journey stages with a source in every row, all six profile-value elements, seven tool roles, and explicit gates for identity, publishing, invitations, messages, personal data, permissions, and unsupported claims. The Review Log must contain the fixed AS-01, Approved Offer, and unsupported '10x' checks; the first two must trace exactly to source, while '10x' must be labelled UNKNOWN and removed.

Write and Boundary-Test the LinkedIn G-C-A-T-E Prompt Contract

LAB 2

You convert the foundation into an operating contract for a recommendation-only LinkedIn assistant. You test a deliberately vague prompt, a grounded profile-rewrite task, missing-evidence behaviour, hostile instructions, voice consistency, and attempts to trigger unauthorised action.

You'll build

C386-linkedin-pack/02-linkedin-agent-contract.md with the G-C-A-T-E contract, output schema, permission matrix, six test cases, observed results, revisions, and release checklist.

Tools: Approved AI assistant, text editor, Lab 1 foundation and supplied current-profile draft

Write and Boundary-Test the LinkedIn G-C-A-T-E Prompt Contract

STEP 1 OF 2

1

Translate the foundation into G-C-A-T-E instructions, an evidence rule, output schema, permission matrix, and explicit stop behaviour.

\$ GOAL + CONTEXT + AUDIENCE + TASK + EVIDENCE + EVALUATION

Write and Boundary-Test the LinkedIn G-C-A-T-E Prompt Contract

STEP 2 OF 2

2

Run adversarial and missing-context tests, revise the contract, and release only when unsupported claims and denied actions are blocked.

\$ TEST -> OBSERVE -> REVISE -> RE-TEST

Write and Boundary-Test the LinkedIn G-C-A-T-E Prompt Contract

Test it

Open 02-linkedin-agent-contract.md. Re-run fixed Test 1 and fixed Tests 4-6 with the fresh-chat protocol. Test 1 must cite supplied evidence, label unknowns, preserve the voice rules, and remain DRAFT; Tests 4-6 must return STOP, avoid the action, and identify the human owner. All six fixed cases must show exact input, expected behaviour, observed behaviour, result, contract version, and any revision.

Recap — Getting Started with Agentic AI for LinkedIn Marketing

1

You can now: LO1, LO2: Define the LinkedIn decision chain, profile value stack, authorised evidence, tool roles, risks, and human-approval gates.

2

You can now: LO1, LO2: Create and test a reusable LinkedIn agent instruction that grounds profile and marketing drafts in approved evidence, voice, permissions, and stop conditions.

TOPIC 02

Content Creation and Personal Branding with AI

Ideas and hooks | posts and articles | content calendar | visuals and carousels | authentic voice

KEY CONCEPTS

Key Concepts — Content Creation and Personal Branding with AI

1

Audience job

Every item answers a professional question, reduces uncertainty, or enables a useful next step.

2

Content pillars

A small set of repeatable themes makes the brand recognisable without becoming repetitive.

3

Hook-to-value

The opening earns attention; the body delivers evidence or insight; the close offers a proportionate next action.

4

Format fit

Short posts, articles, documents, images, and video serve different depth and consumption needs.

5

Series logic

A calendar sequences ideas so one item prepares the audience for the next instead of publishing isolated pieces.

6

Voice fingerprint

Sentence rhythm, vocabulary, stance, examples, and boundaries make a draft sound like the real professional.

KEY CONCEPTS

Key Concepts — Content Creation and Personal Branding with AI (cont.)

7

Creative brief

A visual or carousel starts with one message, a slide-by-slide hierarchy, accessibility, and rights checks.

8

Review before release

Fact, claim, voice, rights, confidentiality, and relevance checks occur before scheduling or publication.

The Audience-Value-Action Content Model

1

Useful LinkedIn content begins with the audience's professional question, not with a format or a request to promote an offer. The content provides one clear value unit such as a framework, interpretation, example, checklist, or informed point of view, then proposes a proportionate action.

2

A post about procurement risk might explain three warning signs and invite readers to compare their own process. The call to action should match readiness: reflect, comment, save, read, follow, or begin a relevant conversation. A strong draft does not force every reader directly into a sales call.

The Audience-Value-Action Content Model — Concept Map

1

Audience question

What is the reader trying to decide or do?

2

Value unit

What insight, example, or tool will help?

3

Proof

Which approved evidence supports the point?

4

Proportionate action

What natural next step fits readiness?

Hooks, Bodies, and Credible Calls to Action


1

A hook creates a relevant information gap without exaggeration. It can name a costly misconception, a surprising but supported observation, a concrete problem, or a useful promise. The body must then pay off that promise with evidence, reasoning, and a practical takeaway.

2

The close should not undermine trust with a generic pitch. It can invite a specific professional response, offer a useful resource, or name the next item in a series. Claims, client outcomes, and quotations must be traceable to approved sources before they appear in any draft.

Hooks, Bodies, and Credible Calls to Action — Concept Map



Choose the Right LinkedIn Format

1

A short post is suited to one focused idea; an article supports sustained reasoning; a document post can turn a framework into a swipeable sequence; an image can make a single relationship memorable; and video can demonstrate a process or point of view. Format follows the audience job.

2

LinkedIn supports text, images, video, polls, events, documents, and articles. Interface details can change, so the workflow should store the content objective and source material independently of a specific button or layout. The human publisher verifies the current interface at release time.

Choose the Right LinkedIn Format — Concept Map

1

Short post

One idea, fast context, discussion.

2

Article

Longer argument, evidence, and depth.

3

Document/carousel

Sequenced framework or tutorial.

4

Image/video

Visual explanation, proof, or demonstration.

Build a Voice Fingerprint

1

Authenticity can be specified. A voice fingerprint describes preferred vocabulary, sentence length, stance, level of formality, recurring examples, claims the author will not make, and words that feel unlike them. Approved writing samples are stronger evidence than vague labels such as professional or engaging.

2

AI should propose alternatives, not erase the author's point of view. The owner compares the draft with the voice fingerprint, rewrites at least one passage, and records why. This review loop gradually improves the instruction set without pretending the model has become the person.

Build a Voice Fingerprint — Concept Map

1

Sound

Rhythm, formality, and preferred vocabulary.

2

Substance

Point of view, examples, and evidence style.

3

Boundaries

Banned clichés, unsupported claims, private topics.

4

Human edit

Owner restores nuance and accountability.

Content Calendar as a Learning Sequence


1

A content calendar is a hypothesis about how ideas compound. It balances content pillars, audience stages, formats, evidence availability, and production capacity. Each item has an objective, owner, source, review status, planned date, and intended measurement.

2

The calendar should preserve slack for timely responses and should not automate publication simply because a date arrived. LinkedIn's native scheduler can hold approved posts, but the final owner still reviews facts, tone, links, media, audience, and timing before scheduling.

Content Calendar as a Learning Sequence — Concept Map



Hands-On Labs — Content Creation and Personal Branding with AI

Lab 3

- Create the LinkedIn Voice, Profile, and Multi-Format Content Kit

Lab 4

- Build the Fourteen-Day Content Calendar and Draft-First Workflow

Create the LinkedIn Voice, Profile, and Multi-Format Content Kit

LAB 3

You build a recognisable content foundation for Northstar Advisory using the same audience value, proof, and action defined in Labs 1 and 2. The AI produces alternatives, while you restore the human point of view and verify every claim, example, quotation, and visual direction.

You'll build

C386-linkedln-pack/03-profile-voice-and-content-kit.md containing the voice fingerprint, reviewed headline and About drafts, four post concepts, one article outline, one six-page carousel brief, evidence maps, and review log.

Tools: Approved AI assistant, text editor, Labs 1-2 outputs and supplied writing samples

Create the LinkedIn Voice, Profile, and Multi-Format Content Kit

STEP 1 OF 2



Derive a voice fingerprint and profile draft from approved writing samples, value, and proof - then let the owner restore nuance.

\$ SOURCE PATTERN -> AI ALTERNATIVES -> HUMAN VOICE

Create the LinkedIn Voice, Profile, and Multi-Format Content Kit

STEP 2 OF 2

2

Build posts, article structure, and carousel as one evidence-linked content kit with fact, voice, privacy, rights, and relevance checks.

\$ HOOK -> VALUE -> PROOF -> PROPORTIONATE ACTION

Create the LinkedIn Voice, Profile, and Multi-Format Content Kit

Test it

Open 03-profile-voice-and-content-kit.md. It must contain eight voice dimensions, three headline options, one About draft, four complete post drafts, an article outline, and a six-page carousel brief. Every factual claim must map to an approved source or be removed. The selected headline, About draft, POST-01 through POST-04, ART-01, and CAR-01 must each contain a documented human edit, and no item may be marked for publication.

Build the Fourteen-Day Content Calendar and Draft-First Workflow

LAB 4

You turn the reviewed content kit into a two-week learning sequence rather than a queue of unrelated posts. You define dependencies, owners, source freshness, review gates, native-scheduling handoff, exception paths, and the signals that will inform the next cycle.

You'll build

C386-linkedin-pack/04-content-calendar-and-workflow.md with a fourteen-day calendar, content dependencies, production board, approval packet, native-scheduling checklist, exception routes, and readiness test.

Tools: Approved AI assistant, text editor or spreadsheet, Labs 1-3 outputs

Build the Fourteen-Day Content Calendar and Draft-First Workflow

STEP 1 OF 2



Sequence reviewed artefacts by audience question and learning dependency, with a source, owner, review date, action, and metric for every row.

\$ CONTENT KIT -> FOURTEEN-DAY LEARNING SEQUENCE

Build the Fourteen-Day Content Calendar and Draft-First Workflow

STEP 2 OF 2

2

Move each item through explicit states and a complete approval packet; hand approved drafts to a human using LinkedIn's native scheduler.

\$ DRAFT -> CHECK -> APPROVE -> HUMAN SCHEDULE -> OBSERVE

Build the Fourteen-Day Content Calendar and Draft-First Workflow

Test it

Open 04-content-calendar-and-workflow.md. The calendar must cover fourteen days, include all six reviewed artefacts plus two engagement days, name a source and owner for every planned item, and use the defined production states. Fixed rows Day 2, Day 7, and Day 12 must trace through a complete review packet to either READY FOR HUMAN SCHEDULING or a specific HOLD reason; the change-evidence table must capture corrections, and no automated publication or live engagement action may exist.

Recap — Content Creation and Personal Branding with AI

1

You can now: LO2, LO3: Apply the prompt contract to create an evidence-grounded profile draft, voice fingerprint, posts, article outline, and carousel brief.

2

You can now: LO3, LO5: Sequence the content kit into a coherent calendar and design a governed workflow from intake through human scheduling and learning.

TOPIC 03

Lead Generation and Outreach with AI Agents

Prospect research | segmentation | personalised invitations and messages | nurture sequences | social-selling etiquette

Key Concepts — Lead Generation and Outreach with AI Agents

1

Authorised evidence

Use supplied, consented, or legitimately accessible information; never scrape profiles or infer sensitive traits.

2

Ideal-customer criteria

Segment by relevant business context, need, fit, timing, and relationship evidence rather than surface popularity.

3

Reason to connect

A message earns attention by naming genuine context and a relevant professional reason.

4

Personalisation

Use only a small number of verified details; do not create false familiarity or pretend an event occurred.

5

Invitation etiquette

Quality and recognisability matter more than volume; excessive or unwanted invitations can restrict an account.

6

Nurture progression

Each touch adds new value and gives the recipient an easy way to decline or disengage.

Key Concepts — Lead Generation and Outreach with AI Agents (cont.)

7

Intent and risk

Questions, objections, complaints, privacy requests, and sensitive topics require different handling and escalation.

8

Human send

The agent drafts and prioritises; the account owner checks the recipient, context, wording, and timing before sending.

Evidence-Based Prospecting Without Scraping

1 Prospect research begins with an ideal-customer hypothesis and authorised evidence. Useful fields include organisation type, role relevance, publicly stated business context, relationship path, observed need, source, date, confidence, and an explicit reason to engage.

2 The workflow must not use scripts, browser extensions, or bots to copy LinkedIn profiles or activity. LinkedIn's User Agreement prohibits scraping and unauthorised automated access. In this course, all prospect work uses synthetic records so learners can practise segmentation without processing a real person's data.

Evidence-Based Prospecting Without Scraping — Concept Map



Segmentation as a Decision Rule


1

A segment is useful when it changes the next action. A simple decision table can combine fit, evidence strength, relationship context, urgency, and risk to produce PRIORITISE, NURTURE, HOLD, or DO NOT CONTACT. The rule should be explainable and reversible.

2

AI may summarise or apply the rule, but it should not invent missing attributes or infer protected or sensitive characteristics. Low confidence should lower the action, not trigger more aggressive data collection.

Segmentation as a Decision Rule — Concept Map



The Relevance Ladder for Outreach


1

Personalisation is not inserting a first name into a generic pitch. Relevance grows from identity, to genuine context, to a specific professional reason, to useful value, and only then to a low-friction next step. Each rung must be true and appropriate to use.

2

If the only available detail is a job title, the message should remain modest. The agent must not fabricate a shared interest, praise content it has not been given, or imply that it reviewed private activity.

The Relevance Ladder for Outreach — Concept Map



Connection, Message, and Nurture Sequence


1

A connection invitation explains who the sender is and why the connection is relevant. After acceptance, a first message should continue the context rather than immediately switch to a hard pitch. Later touches should add distinct value and stop when the recipient declines, does not engage, or the evidence becomes stale.

2

LinkedIn places limits and restrictions on invitations to protect member experience, and advises context and personalisation. Exact product limits can vary by account and time. The workflow therefore never sets a volume target designed to evade restrictions; it emphasises quality, relevance, and owner judgement.

Connection, Message, and Nurture Sequence — Concept Map



Social-Selling Etiquette and Escalation


1

Professional outreach protects the recipient's agency. Messages are accurate, concise, specific, and easy to ignore or decline. The sender does not create false urgency, disguise commercial intent, repeat unanswered requests indefinitely, or move personal data into unapproved tools.

2

Escalation is part of the workflow. Legal questions, privacy requests, complaints, sensitive personal circumstances, high-value commercial commitments, and requests outside the sender's authority go to a named human owner with the original context preserved.

Social-Selling Etiquette and Escalation — Concept Map



Hands-On Labs — Lead Generation and Outreach with AI Agents

Lab 5

- Build the Evidence-Based Prospect Segmentation Queue

Lab 6

- Draft the Personalised Connection and Nurture Sequence

Build the Evidence-Based Prospect Segmentation Queue

LAB 5

You define what a relevant prospect looks like, inspect the supplied synthetic records, and build an explainable decision table. The agent may summarise and apply the rule, but it must preserve provenance, label uncertainty, avoid sensitive inference, and never access LinkedIn or enrich a record automatically.

You'll build

C386-linkedin-pack/05-prospect-segmentation-queue.md with ideal-customer criteria, data dictionary, evidence register, scoring rule, segmented queue, quality checks, exclusions, and human-review log.

Tools: Approved AI assistant, spreadsheet or text editor, supplied synthetic prospect CSV and earlier foundation

Build the Evidence-Based Prospect Segmentation Queue

STEP 1 OF 2

1

Define ideal-customer criteria, authorised fields, hard stops, and an explainable decision order before asking AI to segment anything.

\$ HARD STOP -> EVIDENCE -> FIT -> CONTEXT -> TIMING

Build the Evidence-Based Prospect Segmentation Queue

STEP 2 OF 2

2

Apply the rule to synthetic records, validate every decision, and hand only relevant, verified rows to a human-reviewed message-drafting queue.

\$ PRIORITISE / NURTURE / HOLD / DO NOT CONTACT

Build the Evidence-Based Prospect Segmentation Queue

Test it

Open 05-prospect-segmentation-queue.md. It must define all eleven fields, contain a quality result for every synthetic record, preserve all hard stops, and give every decision a field-level rationale. Apply the written rule manually to the fixed fixtures P-001 (PRIORITISE), P-004 (NURTURE), P-003 (HOLD), and P-007 (DO NOT CONTACT); the same inputs must produce those decisions. HOLD and DO NOT CONTACT must show reason-to-engage as NOT PERMITTED, and no row may contain scraped, enriched, real, or sensitive information.

Draft the Personalised Connection and Nurture Sequence

LAB 6

You turn approved synthetic queue records into a bounded communication sequence. Each draft uses verified context, adds distinct value, avoids false familiarity and pressure, and stops on no response, decline, uncertainty, privacy concern, or relationship risk.

You'll build

C386-linkedin-pack/06-outreach-and-nurture-sequence.md with invitation, welcome, value, and handoff drafts for approved synthetic segments; an evidence map; message checks; stop conditions; and a send-decision log.

Tools: Approved AI assistant, text editor, Lab 5 queue and supplied synthetic message scenarios

Draft the Personalised Connection and Nurture Sequence

STEP 1 OF 2



Build each draft up the relevance ladder using only verified identity and context, a professional reason, useful value, and a low-friction next step.

\$ IDENTITY -> CONTEXT -> REASON -> VALUE -> NEXT STEP

Draft the Personalised Connection and Nurture Sequence

STEP 2 OF 2

2

Route every scenario through etiquette, privacy, stop, and escalation rules; a human verifies the recipient and decides whether a draft is ever sent.

\$ DRAFT RESPONSE / WAIT / STOP / ESCALATE

Draft the Personalised Connection and Nurture Sequence

Test it

Open 06-outreach-and-nurture-sequence.md. Each approved segment must have an invitation, post-acceptance welcome, value follow-up, and optional handoff with a complete evidence map. All seven fixed scenarios MS-01 through MS-07 must have a safe state plus a separate escalation action, reason, and owner. MS-05 must be STOP plus ESCALATE TO PRIVACY OWNER; declined, deletion, disputed, sensitive, and silent cases must not continue automatically. The Send-Decision Log must show no message was sent in the lab.

Recap — Lead Generation and Outreach with AI Agents

1

You can now: LO4: Apply transparent fit, evidence, relationship, timing, and risk rules to authorised synthetic prospect records without scraping or sensitive inference.

2

You can now: LO4, LO5: Create concise, relevant, evidence-linked outreach drafts with etiquette rules, stop conditions, privacy protection, and human sending authority.

TOPIC 04

Automating and Scaling LinkedIn with AI Agents

Agentic content and engagement workflows | scheduling and follow-up | measurement | responsible scale

Key Concepts — Automating and Scaling LinkedIn with AI Agents

1

Workflow states

Intake, validate, draft, review, approve, schedule, observe, and learn are explicit states with owners.

2

Allowed automation

Automate internal preparation and analysis; use LinkedIn's authorised features or approved integrations for external actions.

3

Approval packet

The reviewer sees source evidence, draft, checks, risks, intended audience, and proposed action in one place.

4

Operational trace

Every run records inputs, version, reviewer, decision, action, timestamp, and exception.

5

Metric definition

Impressions, reach, reactions, comments, reposts, saves, sends, profile viewers, and followers answer different questions.

6

Rates need denominators

Engagement and conversion rates are meaningful only when numerator, denominator, scope, and period are explicit.

Key Concepts — Automating and Scaling LinkedIn with AI Agents (cont.)

7

Diagnose before changing

Check data quality, audience mix, content, distribution, and downstream action before blaming the agent.

8

Bounded scale

Change one controlled variable, define a guardrail, owner, observation window, cooldown, and rollback.

A Draft-First LinkedIn Operations Workflow


1

A governed workflow separates internal automation from external platform action. The agent can collect an approved brief, validate required fields, draft content or replies, run checks, and assemble a review packet. A human then approves, revises, or stops the item.

2

External action uses LinkedIn's own features or an organisation-approved integration operating within current terms and permissions. Unauthorised bots that send messages, create engagement, or access the service are outside the workflow boundary.

A Draft-First LinkedIn Operations Workflow — Concept Map



Approval Packet and Exception Path


1

Approval fails when the reviewer has to reconstruct the context. A compact packet includes the objective, intended audience, source links, draft, evaluation results, known unknowns, risk tier, proposed action, and rollback or correction path.

2

Exceptions are normal. Missing evidence returns to intake; a policy concern stops the run; a voice mismatch returns for revision; an urgent complaint escalates to an owner. The workflow records why the path changed so later improvement is based on evidence.

Approval Packet and Exception Path — Concept Map



LinkedIn Measurement Tree


1

Discovery metrics describe exposure; engagement metrics describe interaction; profile activity describes movement toward the professional identity; and business outcomes describe what happened beyond LinkedIn. These layers should be reported separately.

2

LinkedIn notes that member post analytics are estimates and may include the author's own activity. A sound report states the source, extraction date, scope, and limitations. It never converts an estimated impression directly into revenue without an evidence bridge.

LinkedIn Measurement Tree — Concept Map



Metric Formulas and Interpretation

1

A rate is a contract between a numerator and denominator. For a simple synthetic exercise, engagement rate may be defined as total selected interactions divided by impressions. Reply rate may be replies divided by messages sent. The definition must be recorded because tools and teams may include different interactions.

2

Small denominators, short windows, outliers, and changed audience mix can create unstable rates. Report counts beside rates, compare like with like, and avoid declaring a winner from one post or one outreach batch.

Metric Formulas and Interpretation — Concept Map

1

Engagement rate

Selected interactions / impressions x 100.

2

Profile-view rate

Profile viewers from post / impressions x 100.

3

Reply rate

Replies / messages sent x 100.

4

Outcome rate

Qualified outcomes / relevant opportunities x 100.

Responsible Scaling Loop


1

Scaling is a controlled learning loop, not more output. The team validates the data, diagnoses the likely bottleneck, chooses one change, predicts the effect, names a guardrail, obtains approval, observes for a defined window, and then holds, adapts, or rolls back.

2

A bounded recommendation includes the owner, change limit, observation window, cooldown, stop threshold, and rollback. The agent recommends; it does not change publishing frequency, messaging activity, or account behaviour on its own.

Responsible Scaling Loop — Concept Map



Hands-On Labs — Automating and Scaling LinkedIn with AI Agents

Lab 7

- Design the Governed Content and Engagement Operations Workflow

Lab 8

- Diagnose Performance and Propose One Bounded Scaling Experiment

Design the Governed Content and Engagement Operations Workflow

LAB 7

You connect the earlier content calendar, prospect rules, and outreach drafts into one operations workflow. The agent prepares and checks internal artefacts; human owners retain profile identity, publishing, invitations, messages, relationship judgement, and incident correction.

You'll build

C386-linkedin-pack/07-governed-operations-workflow.md containing the workflow map, state machine, role and permission matrix, approval packets, engagement router, exception paths, trace schema, and simulation results.

Tools: Approved AI assistant, text editor or diagram tool, Labs 1-6 outputs and supplied engagement scenarios

Design the Governed Content and Engagement Operations Workflow

STEP 1 OF 2

1

Separate internal preparation from external action with explicit states, role permissions, a complete review packet, and no path that bypasses human approval.

\$ INTAKE -> DRAFT -> CHECK -> HUMAN REVIEW -> HUMAN ACTION

Design the Governed Content and Engagement Operations Workflow

STEP 2 OF 2

2

Route engagement and operational exceptions to a safe state, preserve evidence, name the owner, and log every decision for learning and correction.

\$ RESPOND / WAIT / REVISE / STOP / ESCALATE

Design the Governed Content and Engagement Operations Workflow

Test it

Open 07-governed-operations-workflow.md. Every main and exception state must have an entry rule, owner, allowed next states, and trace fields; no path may bypass human review. All six fixed simulations must finish in their stated expected safe state with a complete record and change-evidence row. READY FOR HUMAN ACTION must appear only after a named human decision; the request to auto-send must finish STOP without executing or simulating an external LinkedIn action.

Diagnose Performance and Propose One Bounded Scaling Experiment

LAB 8

You analyse a supplied synthetic performance export without assuming that reach equals business value. You define each metric, check the data, compare like with like, identify the most plausible bottleneck, and recommend one controlled change that remains inside the governed workflow.

You'll build

C386-linkedin-pack/08-performance-diagnosis-and-improvement.md plus C386-linkedin-pack/08-cleaned-performance-metrics.csv containing definitions, quality log, calculations, funnel diagnosis, alternative explanations, experiment card, approval decision, cooldown, rollback, and course-pack manifest.

Tools: Spreadsheet application, approved AI assistant, supplied synthetic performance CSV and Labs 1-7 outputs

Diagnose Performance and Propose One Bounded Scaling Experiment

STEP 1 OF 2



Define metrics, validate the synthetic export, calculate transparent counts and rates, and separate discovery, engagement, profile, conversation, and business outcomes.

\$ DEFINE -> VALIDATE -> CALCULATE -> INTERPRET

Diagnose Performance and Propose One Bounded Scaling Experiment

STEP 2 OF 2

2

Diagnose alternatives, then propose one controlled variable with an owner, guardrails, observation window, cooldown, stop threshold, and rollback.

\$ HYPOTHESIS -> HUMAN DECISION -> OBSERVE -> HOLD / ADAPT / ROLLBACK

Diagnose Performance and Propose One Bounded Scaling Experiment

✓ Test it

Open 08-performance-diagnosis-and-improvement.md and 08-cleaned-performance-metrics.csv. Verify R-001 engagement 3.83%, R-002 engagement 5.24%, and R-009 reply rate 33.33%; the two R-010 instances must have distinct row_instance_id values and explicit duplicate dispositions, while R-011 impossible reach and R-012 missing denominator must carry explicit dispositions, and blank or zero denominators must return N/A. The diagnosis must state alternatives and limitations, and the experiment must change exactly one variable with an owner, primary metric, two guardrails, observation window, cooldown, stop threshold, and rollback. The course-pack manifest must list outputs 01 through 08.

Recap — Automating and Scaling LinkedIn with AI Agents

1

You can now: LO5: Design a traceable draft-first workflow for content, comments, and follow-ups with explicit states, approvals, exceptions, and authorised external-action boundaries.

2

You can now: LO6: Validate synthetic LinkedIn results, calculate defined metrics, diagnose the decision chain, and produce a human-reviewed improvement with guardrails and rollback.



WRAP-UP

Course Summary & Next Steps

What You Achieved

1

Agent Foundations

Explain how goal-driven AI agents support LinkedIn marketing and define safe tool, evidence, and human-approval boundaries.

2

Profile & Prompts

Optimise a LinkedIn profile and write reusable prompts that preserve an authentic professional voice.

3

Content System

Create evidence-grounded LinkedIn posts, articles, visual briefs, carousels, and a coherent content calendar.

4

Prospecting

Research and segment prospects from authorised information, then draft relevant connection and nurture messages for human review.

5

Governed Workflow

Design a draft-first content, engagement, and follow-up workflow that respects LinkedIn rules and professional etiquette.

6

Measure & Improve

Define LinkedIn performance metrics, diagnose synthetic results, and recommend one bounded improvement with a traceable decision record.

Continue Responsibly

1

Keep the agent recommendation-only until evidence quality and review consistency are demonstrated.

2

Verify current LinkedIn help and terms before any operational change.

3

Measure the complete decision chain, not just impressions or activity volume.

4

Use small controlled changes with an owner, guardrail, observation window, and rollback.

Keep Practising After the Course

1

Re-run the labs on your own machine — repetition is what makes the skills stick.

2

Apply the techniques to a real process or project in your own organisation.

3

Your course material stays available at <https://lms-tms.tertiaryinfotech.com/>.

4

Explore the follow-on courses at www.tertiarycourses.com.sg.

C386 | KEEP THE HUMAN ACCOUNTABLE

Thank You!

You can now design a useful, evidence-grounded LinkedIn marketing workflow in which AI accelerates preparation and people retain responsibility for identity, relationships, and external action.